

# Fast Convergence PageRank in Hadoop

Team Members:

Karthik Bellur(kvb26)

Sumukh SP(ss3345)

Prashanth Basappa(pb476)

# Problem

Compute PageRank for a large scale web graph (685230 nodes and 7600595 edges) using a fast convergence using Blocked PageRanks

Idea is to reduce the I/O cost of a MapReduce job by doing nontrivial computation in the reduce step and increase the performance by reducing the number of passes required to achieve convergence.

# Input Data

We pre-process the raw input data from edges.txt to create filtered\_edges.txt.

We then perform some additional processing on the filtered\_edges.txt file in order to bring it to a format which is convenient for the Mapper process to read.

## Mapper Input File Format:

<NodeID PageRank n D1 D2 D3 ..... Dn>

NodeID = Node ID of a source node

PageRank = Page Rank of the source node

D1 D2 D3 ... Dn = List of destination nodes

```
double fromNetID = 0.5433;           // using netid ss3345
double rejectMin = 0.9 * fromNetID;  // =0.48897
double rejectLimit = rejectMin + 0.01; // =0.49897

boolean selectInputLine(double x) {
    return ( ((x >= rejectMin) && (x < rejectLimit)) ? false : true );
}
```

# Node-by-node Computation of PageRank

## SimplePageRankMaster

- Bootstrapper for the MapReduce process
- Set up the Mapper & Reducer jobs; configure the input/output paths
- Create a MapReduce counter to store “residual”
- Output the residual (so we can conclude that PageRank is still far from convergence)

## SimplePageRankMapper

### Input:

<NodeID PageRank n D1 D2 D3 ..... Dn>

NodeID = Node ID of a source node

PageRank (PR) = Page Rank of the source node

D1 D2 D3 ... Dn = List of destination nodes

n = Total number of nodes in the graph

### Output:

For each node in <D1, D2, D3, ... Dn>

emit( NodeID, PageRank/n )

For each node in the graph

emit( NodeID, “PR\_” + <NodeID PageRank n D1 D2 D3 ..... Dn> )

# Node-by-node Computation of PageRank

## SimplePageRankReducer:

### Input:

<NodeID, Inbound PageRank>

<NodeID, “PR\_” + <NodeID PageRank n D1 D2 D3 ..... Dn>>

### Output:

<NodeID, <NodeID NewPageRank n D1 D2 D3 ..... Dn>>

$\text{NewPageRank (NPR)} = (1-d) / N + \text{Inbound PageRank} * d$

The residual PageRank is calculated by the Reducer as shown below and added to a MapReduce counter.

**Residual PageRank** =  $\text{Math.abs}(\text{NewPageRank} - \text{PageRank}) / \text{NewPageRank}$

The average residual PageRank over all the nodes in the graph is calculated as shown below:

**Avg Residual PageRank** =  $\text{Residual PageRank} / n$ ;

# Results

Iteration 1, Avg Error: 2.3390409059731767  
Iteration 2, Avg Error: 0.3228536403835209  
Iteration 3, Avg Error: 0.19199830713774937  
Iteration 4, Avg Error: 0.09401514819841512  
Iteration 5, Avg Error: 0.06275703048611415  
Iteration 6, Avg Error: 0.03391270084497176  
Iteration 7, Avg Error: 0.027217138770923632  
Iteration 8, Avg Error: 0.0165739970520847  
Iteration 9, Avg Error: 0.01436597930621835  
Iteration 10, Avg Error: 0.009770442041358376  
Iteration 11, Avg Error: 0.008454095705091721  
Iteration 12, Avg Error: 0.006145381842592998

Iteration 13, Avg Error: 0.005370459553726486  
Iteration 14, Avg Error: 0.00402346657326737  
Iteration 15, Avg Error: 0.0035068517140230287  
Iteration 16, Avg Error: 0.002715876421055704  
Iteration 17, Avg Error: 0.0023408198707003487  
Iteration 18, Avg Error: 0.0018533922916393036  
Iteration 19, Avg Error: 0.0015907067699896386  
Iteration 20, Avg Error: 0.0012754841440100403  
Iteration 21, Avg Error: 0.00109014491484611  
Iteration 22, Avg Error: 8.858339535630373E-4

# Blocked Computation of PageRank

## Node to Block Mapping:

The MapReduce master process, on startup, reads from the blocks.txt file and creates a list consisting of the prefix sums of the values contained in the blocks.txt file. A binary search of this list would then give us the corresponding Block for the Node under consideration in  $O(\log(\text{numberOfBlocks}))$  time.

## BlockedPageRankMapper

Input:

<NodeID PageRank n D1 D2 D3 ..... Dn>

NodeID = Node ID of a source node

PageRank (PR) = Page Rank of the source node

D1 D2 D3 ... Dn = List of destination nodes

n = Total number of nodes in the graph

# Blocked Computation of PageRank

## Output:

For each node in the graph

emit( BlockID, "PR " + <NodeID PageRank n D1 D2 D3 ..... Dn> )

For each node in <D1, D2, D3, ... Dn>

If( blockIdOfNode( Di ) == blockIdOfNode( NodeID ) )

emit( blockIdOfNode( Di ), "BE " + NodeID + " " + Di )

else

emit( blockIdOfNode( Di ), "BC " + NodeID + " " + Di + " " + PR / n )



# Blocked Computation of PageRank

**BlockedPageRankReducer:**

**do until convergence:**

For each Value for block b

BE <- Set of edges in the same block b

BC <- Set of incoming edges to block b

For each node v in block b

For each node u: edge(u,v) is in b

sum incoming page rank for v

For node u : edge(u,v) is an incoming edge for a different block b'

sum incoming page rank for v

$\text{newPageRank}(v) = (1-d)/n + d * (\text{sum of incoming page ranks to } v)$

**Residual PageRank** =  $\text{Math.abs}(\text{NewPageRank} - \text{PageRank}) / \text{NewPageRank}$

emit (v , new PageRank(v), n , D1,D2,...,Dn)

**Convergence Condition:**

Average Residuals / N <= 0.01

# Results

Iteration 1, Avg Error: 2.815142360959094, Avg iterations per block = 17.58823529411765

Iteration 2, Avg Error: 0.03704449600863943, Avg iterations per block = 7.132352941176471

Iteration 3, Avg Error: 0.02345781708331509, Avg iterations per block = 5.955882352941177

Iteration 4, Avg Error: 0.009277177006260672, Avg iterations per block = 3.911764705882353

Iteration 5, Avg Error: 0.003460152065729755, Avg iterations per block = 2.411764705882353

Iteration 6, Avg Error: 0.001113494738992747, Avg iterations per block = 1.5294117647058822

Iteration 7, Avg Error: 6.041766997942297E-4, Avg iterations per block = 1.1323529411764706

# Gauss Siedel PageRank

Nodes in the graph are sorted.

New PageRank values are used in successive computation of PageRank values in the same iteration.

## Results:

Iteration 1, Avg Error: 2.815524714329495, Avg iterations per block = 10.088235294117647

Iteration 2, Avg Error: 0.038112750463348076, Avg iterations per block = 5.073529411764706

Iteration 3, Avg Error: 0.02474789486741678, Avg iterations per block = 4.470588235294118

Iteration 4, Avg Error: 0.010522014506078252, Avg iterations per block = 3.2205882352941178

Iteration 5, Avg Error: 0.004513812880346745, Avg iterations per block = 2.3676470588235294

Iteration 6, Avg Error: 0.001625731506209593, Avg iterations per block = 1.6176470588235294

Iteration 7, Avg Error: 7.413569166557215E-4, Avg iterations per block = 1.2647058823529411

# Random Blocked PageRank Computation

Randomly assign nodes to blocks instead of using the METIS graph partition

We used a simple function:  $(\text{nodeId} * 541) \% 68$ ;

## Results:

Iteration 1, Avg Error: 2.33966405440509  
Iteration 2, Avg Error: 0.3223165944281482  
Iteration 3, Avg Error: 0.19119565693270874  
Iteration 4, Avg Error: 0.0934357806867767  
Iteration 5, Avg Error: 0.06209010113392584  
Iteration 6, Avg Error: 0.0335172131984881  
Iteration 7, Avg Error: 0.02690045678093487  
Iteration 8, Avg Error: 0.016401792099003255  
Iteration 9, Avg Error: 0.014182099441063585  
Iteration 10, Avg Error: 0.009666827196707676

# EMR Deployment

▼ Steps

Add stepClone step

Steps

View all interactive jobs | View all jobs

Filter: All stepsFilter steps ...

5 steps (all loaded)

	ID	Name	Status	Start time (UTC-4) ▼	Elapsed time	Log files
▶	s-3SBNCXUV3HPK	Run Hue	Completed	2015-05-05 06:50 (UTC-4)	1 minute	<a href="#">View logs</a>
▼	s-lMR5906QS95T	SimplePageRank	Running	2015-05-05 06:52 (UTC-4)	13 minutes	<a href="#">View logs</a>
<div>JAR location: s3://edu-cornell-cs-cs5300s15-ss3345/job/page_rank.jar</div> <div>Main class: None</div> <div>Arguments: com.simple.SimplePageRankMaster s3://edu-cornell-cs-cs5300s15-ss3345/input/ s3://edu-cornell-cs-cs5300s15-ss3345/output</div> <div>Action on failure: Terminate cluster</div>						
▶	s-29HY0YXPOFADO	Setup pig	Completed	2015-05-05 06:50 (UTC-4)	6 seconds	<a href="#">View logs</a>
▶	s-2GPSLYW144EPI	Setup hive	Completed	2015-05-05 06:49 (UTC-4)	24 seconds	<a href="#">View logs</a>
▶	s-2M5LK7D4L6702	Setup hadoop debugging	Completed	2015-05-05 06:49 (UTC-4)	5 seconds	<a href="#">View logs</a>



Thank you!